Planbureau voor de Leefomgeving

PEARL in de cloud

Trond Grytli Husby

Netherlands Environmental Assessment Agency (PBL)

trond.husby@pbl.nl

Den Haag 27 June, 2019

Regional population projections of PBL and CBS

- Every three years the PBL and CBS make the Regional (and National) Population Projections
- Projections of population, households and demographic events in Dutch municipalities until 2050
- Previous edition was in 2016, next in September 2019

Source: https://img.clipartxtras.com/03843abee543e741be870c4ada22c760_free-to-use-public-domain-cauldron-clip-art-witch-cauldron-clipart_500-500.png

Cohort-component model

$$P_{t+1} = P_t + B - X + I - E + A - D$$

P_t: population in t

B: births in the interval $(t, t+1)$

X: deaths in the interval $(t, t+1)$

I: immigration in the interval $(t, t+1)$

E: emigration in the interval $(t, t+1)$

A: arrivals in the interval $(t, t+1)$

D: departures in the interval $(t, t+1)$

- PEARL is a book-keeping system, per definition no uncertainty (unless coding errors)
- But considerable uncertainty in the (exogenous) time-path of growth components
- We can obtain a consistent estimate of the projection uncertainty with Monte Carlo Simulation
- An attempt to carry out a fully stochastic projection stranded on computing power

Azure case 1: estimation of short-distance migration model [1]

- Short-distance migration in PEARL governed by estimated parameters
- Gravity model, estimated with geographically weighted Poisson regression (maximum likelihood)

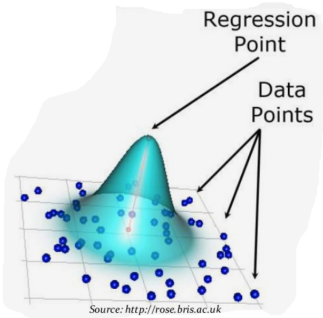

[1] https://github.com/trondhusby/GWmodel

Results

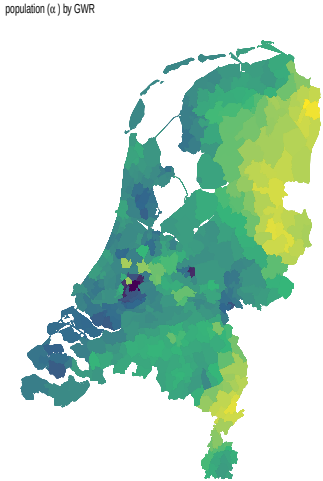
population (α) by GWR

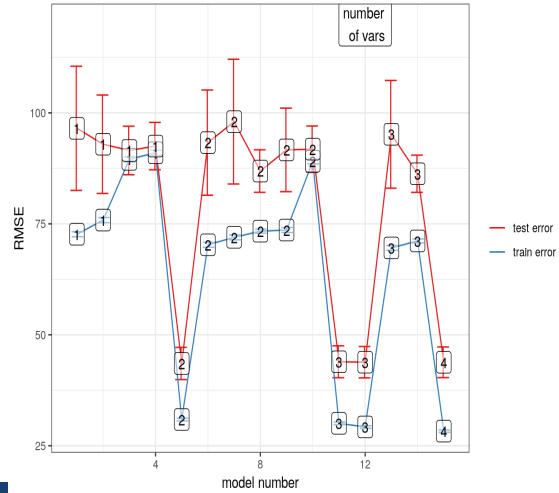
number
of vars
RMSE
100
75
50
25
4
8
12
model number
test error
train error

Why Azure?

- GWR: one model per regression point (16000) times number of MLE iterations. The entire run typically $\tilde{4}$ hours
- Kernel type selection: 4 different kernel types, OLS and MLE (8 models)
- Variable selection: 10-fold cross validation of MLE for all permutations of 4 variables (15 models times 10)
- Bandwidth selection: 10-fold cross validation of MLE for four models from previous step for 7 bandwiths (28 models times 10)

Azure and R

- R package *rAzureBatch*: interaction with the Azure Batch service's REST API
- R package *doAzureParallel*: parallel backend for the *foreach* package.
- `out <- foreach(i=1:n, .options) %dopar% {func(i)}`
- 12 - 15 F64s v2 / F72s v2 Linux VMs in parallel
- Copy of the user's existing R environment stored in Azure Storage. Each VM fetches and loads the R environment. The VM runs the R code inside each iteration of the foreach loop under the loaded R environment. Once the code is finished, results are pushed back to Azure Storage, and a merge task is used to aggregate the results.
- Note: high-level parallelisation in R not always efficient. Need to check whether low-level parallelisation via multicore BLAS/LAPACK is in place

Run-time access to node via ssh

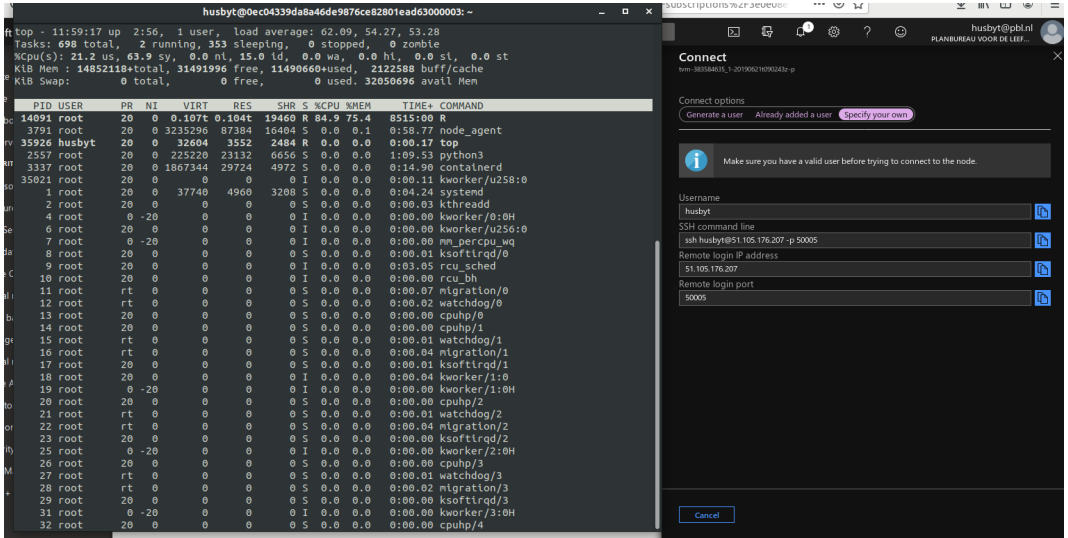

Run-time access to log via Batch Explorer

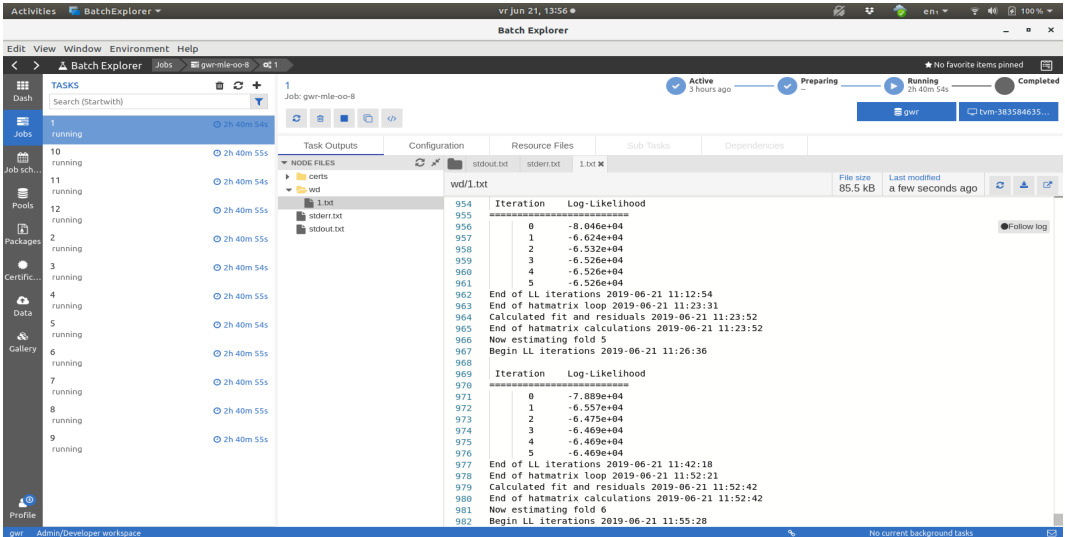

Azure case 2: Monte Carlo simulation with PEARL

- Planned net new dwellings per municipality per year (zachte plannen) have been obtained through a questionnaire among provincial- and municipality-level planners
- 60 municipalities responded, the rest have been imputed
- A major source of uncertainty: do the planners (or we) know building construction until 2050?
- PoC with Azure: run 100 simulations where stochastic construction plans are created with a simple method

Stochastic construction plans

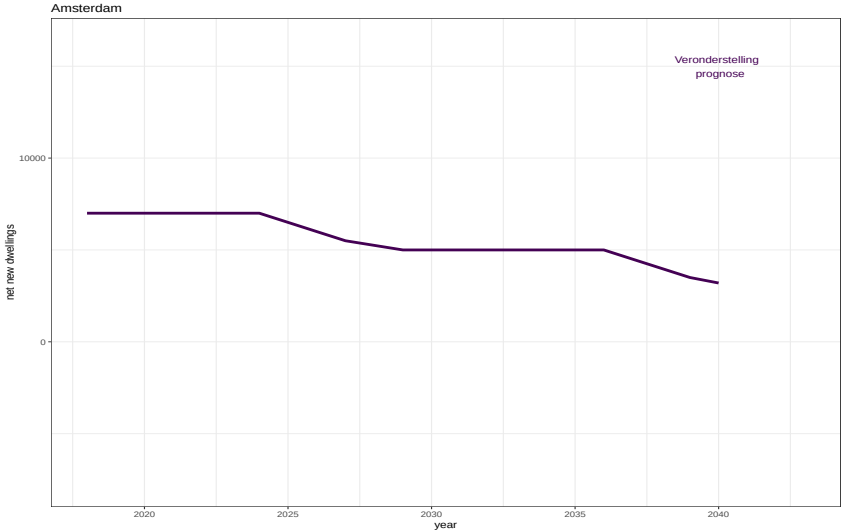

Stochastic construction plans

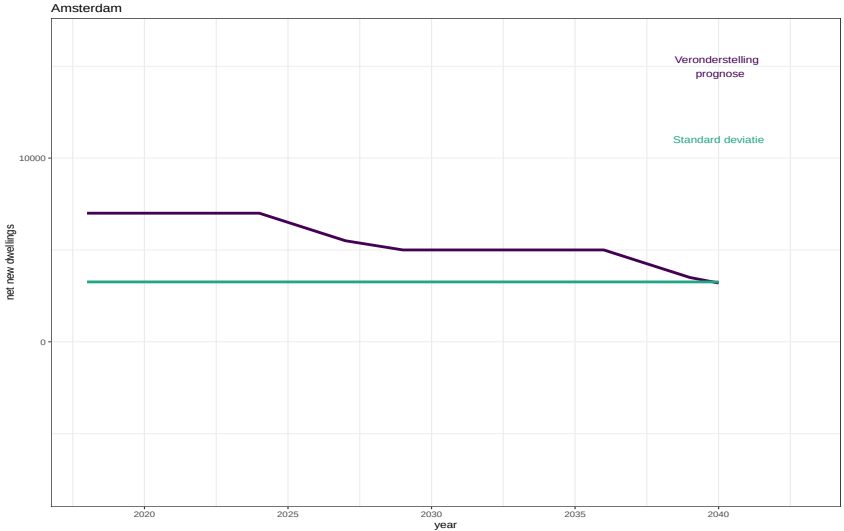

Stochastic construction plans

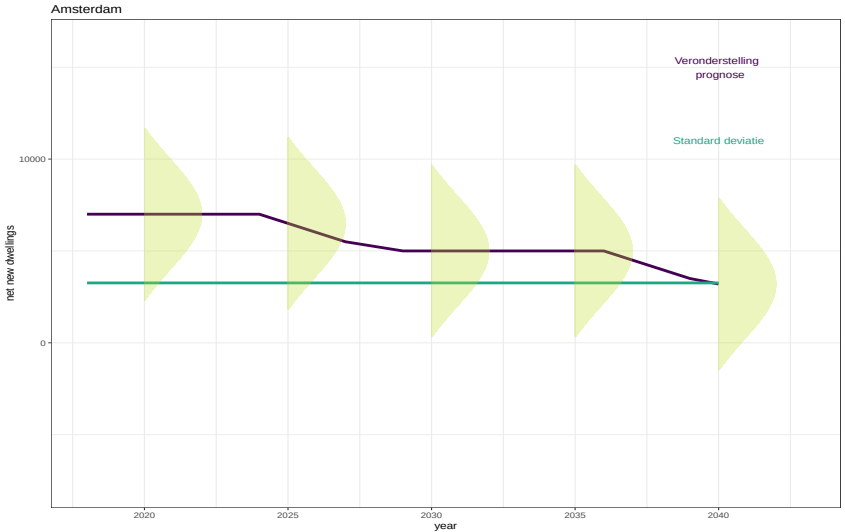

Stochastic construction plans

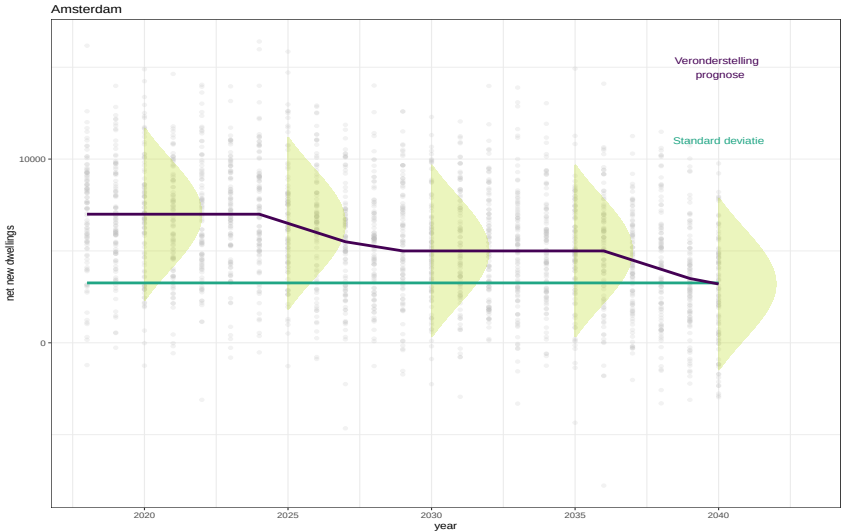

PEARL in de cloud

- PEARL: vb.net application; settings with .ini; and ASCII input files. Run time $\tilde{4}0$ min
- Connection with Azure using Ruben's python script from the PBL Onderzoeksomgeving
- 100 D1 v1 Windows VMS
  - Folders: Application, InputFiles, OutputFiles
  - *config.py*: settings
  - *starttask.cmd*: unzip files to shared dir
  - *task.cmd*: on each VM, copy files from shared dir, run model, create zip files of output
  - Output files pushed back to Storage upon completion of each task
- Separate static input files from simulation inputs into separate folders / zip files
- Download files from storage

PEARL in de cloud

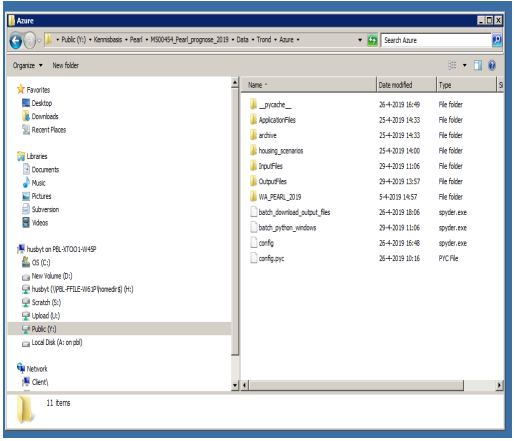

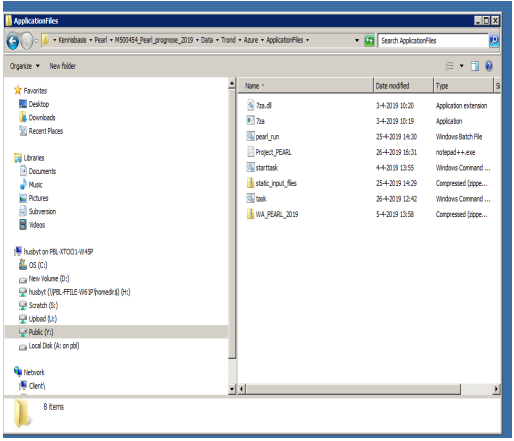

PEARL in de cloud

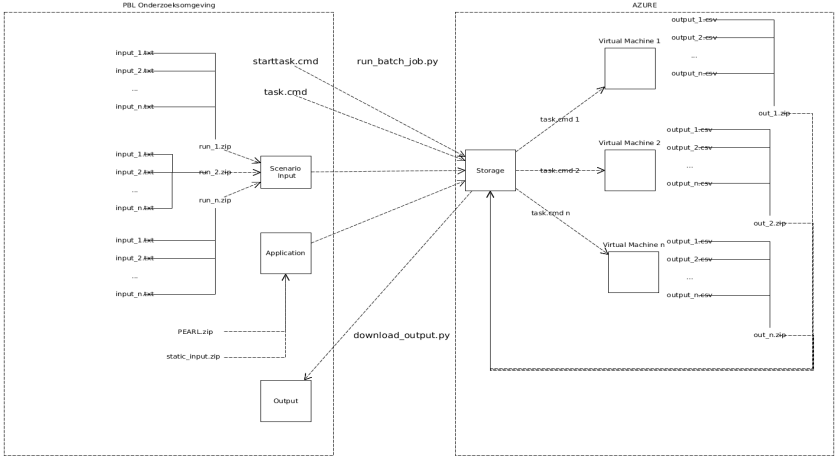

PEARL in de cloud

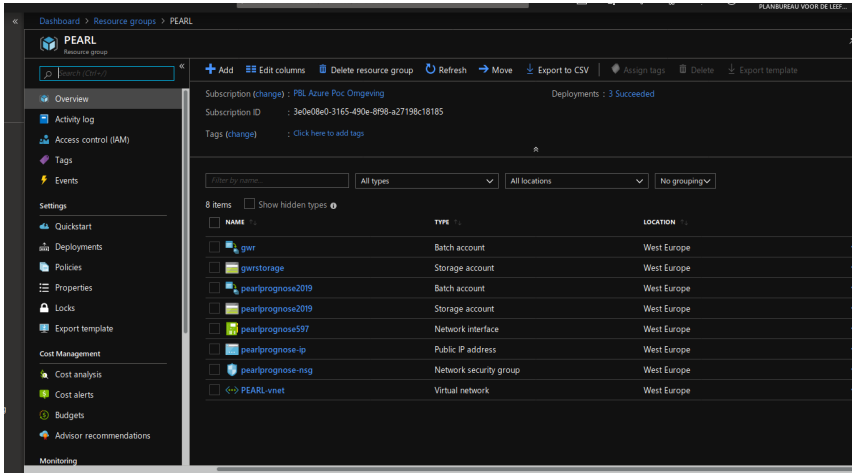

Results

Figure: Short distance arrivals

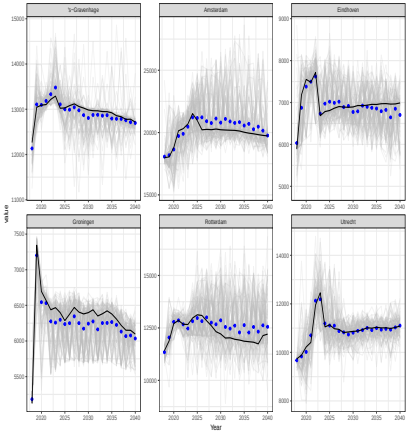

Figure: Population

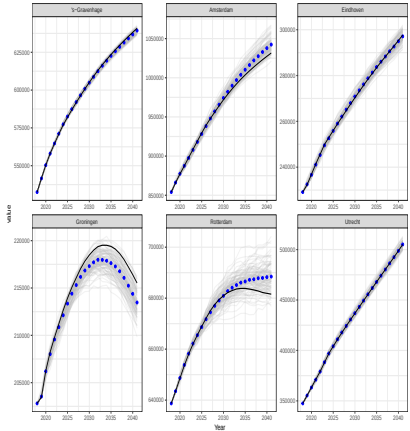

Conclusions

- Two use cases
  1. Azure opens new possibilities for a computationally heavy model (systematic check of out-of-sample performance)
  2. With Azure we can do fully stochastic population projection (current method is *quasi-stochastic*, with a limited amount of runs)
- Running Azure from R is not so straight forward. Some of the usual suspects (installation of packages; dependencies...) caused problems. Working with Docker images a possible solution
- Ruben's python script does the job but could use some extensions: number of tasks is defined by zip-files in the input folder, no autoscaling options...
- Last but not least: computing power should not be a substitute for efficient code